
Word and sentence embeddings

Sentence Classification with LSTMs/ConvNets

Mario Ynocente Castro
Master MVA
ENS Cachan / Ecole Polytechnique
mario.ynocente-castro@polytechnique.edu

1 Word and sentence embeddings with word2vec

1.1 Loading models

1. What is the total number of raw words found in the corpus?
17,005,207 words.
2. What is the number of words retained in the word2vec vocabulary (with default min_count = 5)?
71290 words.

1.2 Exploring the embedding space

1. What is the similarity between (apple and mac), between (apple and peach), between (banana and peach)? In your opinion, why are you asked about the three previous examples?

- Similarity between apple and mac: 0.567861632452
- Similarity between apple and peach: 0.178399832237
- Similarity between banana and peach: 0.688715470006

This illustrates that in this embedding space apple and mac are more similar than what apple and peach are even though both are fruits, in the case of banana and peach we do get a high similarity. This may happen because in the corpus the word apple is used more often as referring to the mark and less referring to the fruit, so it is more common to find it in the same context that the word mac.

2. What is the closest word to the word 'difficult', for 'model' and 'model_phrase'. Comment about the difference between model and model_phrase. Find the three phrases that are closest to the word 'clinton'.

- Closest word to difficult for model: easy
- Closest word to difficult for model_phrase: very_difficult
- Difference between model and model_phrase: model_phrase was trained after pre-processing the tokens of the sentences in the corpus to group the ones that commonly occur together in phrases, model in contrast was trained directly on the individual tokens.
- Three phrases that are closest to clinton: bush, reagan and gore are the top and consist of only one token, the top consisting of two tokens are bill_clinton, w_bush and al_gore.

3. Find the closest word to the vector "vect(france) - vect(germany) + vect(berlin)" and report its similarity measure.

The closest word is 'paris' with similarity measure of 0.757699728012085

4. Explore the embedding space using these functions and report some interesting behaviour (of your choice)
 - We can further confirm that the word apple is more frequently used (on the corpus) as related to the mark by checking which are its most similar words, which gives us: macintosh, atari, amiga, intel, ibm, pc
 - We can make analogies like:
 - science - scientist + mathematician \approx mathematics
 - science - scientist + physicist \approx physics
 - science - scientist + philosopher \approx philosophy
 - science - scientist + astronomer \approx astronomy
 - science - scientist + biologist \approx humanities

1.3 Sentence embeddings

1. Report the closest sentence to the sentence with idx "777", and their similarity score.
"gymnasts get ready for a competition ." with score 0.902949842134
2. Report the 5 closest sentences to the sentence with idx "777", and the associated similarity scores.

Sentence	Similarity score
gymnasts get ready for a competition .	0.902949842134
a woman is getting ready to perform a song for the audience .	0.890097422822
a runner in a competition want to go to the finish line .	0.855536495002
men working to build a fence for customers .	0.851471676783
a man prepares to give a speech for a television audience .	0.849476121272

1.4 IDF weighted sentence embeddings

1. Report the IDF score of the word "the", the word "a", and the word "clinton".
The word "the" has a score of 0.867762351618
The word "a" has a score of 0.473266274881
The word "clinton" doesn't have an IDF score since it's not present in the data.
2. Report the closest sentence to sentence with idx 777.
The closest sentence is "gymnasts get ready for a competition ." with a score of 0.897237646962.

2 Simple LSTM for Sequence Classification

1. What is the (minibatch) shape of:
 - the input of the embedding layer: 32×80
 - the input of the LSTM layer: $32 \times 80 \times 32$
 - the output of the LSTM layer 32×64
2. Report the number of parameters of the model with the standard set of hyper-parameters. Report also the number of sentences in the training set. In standard statistics, a rule of thumb is to have less parameters than samples in your dataset. How do you think its possible to train a model that has so many parameters compared to this number of samples?
 - Number of hyper-parameters: 480000
 - Number of sentences in the training set: 35000
 - We are adding some regularization to the model, in this particular case we use dropout, which can be seen as training multiple similar models with some constraints of weight sharing.

3. For a single sentence, the LSTM has states h_1, \dots, h_T where T is the number of words in the sentence. The sentence embeddings that is fed to the classifier is thus computed as $f(h_1, \dots, h_T)$. What is the exact form of $f(h_1, \dots, h_T)$ used in the python script?

Keras implements LSTM according to the following recurrent formulas:

$$\begin{aligned} f_t &= \sigma(W_f x_t + U_f h_{t-1} + b_f), i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i) \\ C_t &= f_t C_{t-1} + i_t \tanh(W_c x_t + U_c h_{t-1} + b_c) \\ o_t &= \sigma(W_o x_t + U_o h_{t-1} + b_o), h_t = o_t \tanh(C_t) \end{aligned}$$

If the flag `return_sequences` is set to `False` (which is the default) then $f(h_1, \dots, h_T) = h_T$, if it is `True` then $f(h_1, \dots, h_T) = [h_1; \dots; h_T]$.

4. Plot the evolution of the train and valid accuracy per epoch, and write the test errors that you obtain.




Figure 1: Evolution of the loss and accuracy with no dropout.




Figure 2: Evolution of the loss and accuracy with dropout.

- Results without dropout: loss = 0.656337736861, error = 18.0866666698%
- Results with dropout: loss = 0.377996407843, error = 16.60666666699998%

5. Explain what is the difference between SGD and Adam.

The difference is that Adam uses an exponentially decaying average of the past gradients and also divides this by another similarly calculated average, but this time of the sum of squares of the terms in the gradient, which has the effect of decreasing the learning rate of the parameters that have been updated the most.

3 Simple ConvNet for Sequence Classification

1. Report the results (test loss and test error) that you obtain.

Loss = 0.361414234543

Error = 16.3333333365%




Figure 3: Evolution of the loss and accuracy with 1D Convolution.

2. What is the input and output shape of Convolution1D?

Input's shape: $32 \times 80 \times 16$

Output's shape: $32 \times 78 \times 250$

3. Build a model where on top of the convolution, you have an LSTM. It means that the input of the LSTM will be the output of your ConvNet. Run the model with the best parameters you find. Report your best results.

Loss = 0.338560722351

Error = 14.67333333299998%




Figure 4: Evolution of the loss and accuracy with LSTM on top of 1D Convolution.